


3.1 Operators based on first derivative

Let's suppose an image processing technique that pursues the finding of those pixels in images that exhibit brightness variations, that is, their intensity differ from that of their neighbor pixels. Such pixels are often called **edges**, so this technique receives the name of **edge detection**. Edge detection has many applications in computer vision tasks like image segmentation or data extraction. We will explore some of them later in this book.

As commented, edges can be defined as transitions between image regions that have different gray levels (intensities). In this way, the unidimensional, continuous model of an ideal edge is:



That is, an edge is defined by three parameters: its change in intensity, the angle of its slope, and its midpoint. However, in real images that model is not exactly followed by edges since images are discrete, and they are corrupted by noise.

Moreover, **the nature of edges may be diverse:**

- occlusion borders,
- different orientation of surfaces,
- different reflectance properties,
- different textures,
- illumination effects: shadows, reflections, etc.

Error types related to edge detection

Finding edges properly is not a straightforward task, as there exist different errors that can appear when applying edge detection techniques:

- **Detection error.** A good detector exhibits a low ratio of false negative and false positive, that is:
 - False negatives: Existing edges that are not detected.
 - False positives: Detected objects that are not real.
- **Localization error.** Edges are detected, but they are not at the real, exact position.
- **Multiple response.** Multiple detections are raised for the same edge (the edge is thick).

The following figure illustrates such errors.





True
edge



Poor robustness
to noise



Poor
localization




Too many
responses

Thereby, when designing a good edge detector, the goal is to achieve low detection and localization errors, as well as to avoid multiple responses.

Operators based on first derivative (gradient)

In the upcoming chapters, we are going to investigate and implement different edge detection methods. All of them are based on our dear convolution operation, having their own pros and cons.

Concretely, in this notebook we will cover **first-derivative** based operators, which try to detect borders by looking at abrupt intensity differences in neighbor pixels. In the image below we can see two functions $f(x)$ (first row) and how their derivatives (second row) reach their maximum values at the points where the functions' values change more abruptly (around x_0).



If we are dealing with a **two-dimensional** continuous function $f(x, y)$, its derivative is a **vector (gradient)** defined as:

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial}{\partial x} f(x, y) \\ \frac{\partial}{\partial y} f(x, y) \end{bmatrix} = \begin{bmatrix} f_x(x, y) \\ f_y(x, y) \end{bmatrix}$$


which points at the *direction* of maximum (positive) variation of $f(x, y)$:

$$\alpha(x, y) = \arctan\left(\frac{f_y(x, y)}{f_x(x, y)}\right)$$

and has a *module* proportional to the strength of this variation:

$$|\nabla f(x, y)| = \sqrt{(f_x(x, y))^2 + (f_y(x, y))^2} \approx |f_x(x, y)| + |f_y(x, y)|$$

The image below shows examples of gradient vectors:



Concretely, the techniques based on the first derivative explored here are:

- Discrete approximations of a **gradient operator** (Sobel, Prewitt, Roberts, etc., [Section 3.1.1](#)).

Problem context - Edge detection for medical images

Edge detection in medical images is of capital importance for the diagnosis of different diseases (e.g., the detection of tumor cells) in human organs such as lungs and prostates, becoming an essential pre-processing step in medical image segmentation.



In this context, *Hospital Clínico*, a very busy hospital in Málaga, is asking local engineering students to join their research team. They are looking for a person with knowledge in image processing and, in order to ensure it, they have published 3 medical images: `medical_1.jpg`, `medical_2.jpg` and `medical_3.jpg`. They have asked us to perform accurate edge detection in the three images, as well as to provide an explanation of how it has been made.

```

In [28]: import numpy as np
         from scipy import signal
         import cv2
         import matplotlib.pyplot as plt
         import matplotlib
         from ipywidgets import interactive, fixed, widgets
         matplotlib.rcParams['figure.figsize'] = (15.0, 15.0)

         images_path = './images/'

         def sobel_h_55():
             return np.array([
                 [1,2,0,-2,-1],
                 [2,3,0,-3,-2],
                 [3,5,0,-5,-3],
                 [2,3,0,-3,-2],
                 [1,2,0,-2,-1],
             ])*1/9

         def sobel_v_55():
             return np.transpose(sobel_h_55())*-1

         def general_h_kernel(k):
             return np.array([
                 [1,0,-1],
                 [k,0,-k],
                 [1,0,-1]
             ])*1/(2+k)

         def general_v_kernel(k):
             return np.transpose(general_h_kernel(k))*-1

```

To face this challenge, we are going to use plenty edge detection methods, which will be tested and compared in order to determine the best option.

ASSIGNMENT 1: Taking a look at images

First, **display the provided images** to get an idea about what we are dealing with.

Note: As most medical images does not provide color information, we are going to use border detection in grayscale images.

Tip: Different approaches can be followed for edge detection in color images, like converting to YCrCb color space (appendix 2), or detecting edges on each RGB channel.

```

In [29]: # ASSIGNMENT 1
         # Display the provided images in a 1x3 plot to see what are we deal.
         # Write your code here!

         # Read the images
         medical_1 = cv2.imread(images_path + 'medical_1.jpg', 0)
         medical_2 = cv2.imread(images_path + 'medical_2.jpg', 0)
         medical_3 = cv2.imread(images_path + 'medical_3.jpg', 0)


         # And show them
         plt.subplot(131)

```

```
plt.imshow(medical_1, cmap='gray')
plt.title('Medical 1')

plt.subplot(132)
plt.imshow(medical_2, cmap='gray')
plt.title('Medical 2')

plt.subplot(133)
plt.imshow(medical_3, cmap='gray')
plt.title('Medical 3')
plt.show()
```



3.1.1 Discrete approximations of a gradient operator

The first bunch of methods that we are going to explore carry out a **discrete approximation of a gradient operator** based on the differences between gray (intensity) levels. For example, in order to obtain the derivative in the rows' direction, we could apply:

- Backward difference of pixels along a row:

$$f_x(x, y) \approx G_R(i, j) = [F(i, j) - F(i - 1, j)]/T$$

0	0	0
0	1	-1
0	0	0

- Symmetric difference of pixels along a row:

$$f_x(x, y) \approx G_R(i, j) = [F(i + 1, j) - F(i - 1, j)]/2T$$

0	0	0
1	0	-1
0	0	0

These approximations are typically implemented through the convolution of the image with a pair of templates H_C (for columns, detecting vertical edges) and H_R (for rows, detecting horizontal ones), that is:

$$G_R(i, j) = F(i, j) \otimes H_R(i, j)$$

$$G_C(i, j) = F(i, j) \otimes H_C(i, j)$$

Perhaps the most popular operator doing this is such of **Sobel**, although there are many of them that provide acceptable results. These operators use the aforementioned two kernels (typically of size 3×3 or 5×5) which are convolved with the original image to calculate approximations of the derivatives.

These are some examples (first column: operator name; second one: H_R ; third column:

H_C :

Roberts

0	0	0
0	0	1
0	-1	0

-1	0	0
0	1	0
0	0	0

Prewitt

$\frac{1}{3}$

1	0	-1
1	0	-1
1	0	-1

$\frac{1}{3}$

-1	-1	-1
0	0	0
1	1	1

Sobel

$\frac{1}{4}$

1	0	-1
2	0	-2
1	0	-1

$\frac{1}{4}$

-1	-2	-1
0	0	0
1	2	1

Frei-Chen

$\frac{1}{2+\sqrt{2}}$

	1	0		-1
$\sqrt{2}$		0	$\sqrt{2}$	-
	1	0		-1

$\frac{1}{2+\sqrt{2}}$

-1	$\sqrt{2}$	-	-1
0		0	0
1	$\sqrt{2}$		1

In general

$\frac{1}{2+K}$

	1	0		-1
K		0	K	-
	1	0		-1


$\frac{1}{2+K}$


-1	K	-	-1
0		0	0
1	K		1

At this point we know how to perform a discrete approximation of a gradient operator through the application of a convolution operation with two different kernels, that is:

$$\nabla F(x, y) = \begin{bmatrix} F \otimes H_C \\ F \otimes H_R \end{bmatrix}$$

But, how could we use the output of those computations to detect edges? The following figure clarifies that!





Kernel sizes

As discussed, kernels can be of different size, and that size directly affects the quality of the detection and the localization (e.g. Sobel 3×3 or 5×5):

- Small template:
 - more precise localization (good localization).
 - more affected by noise (likely produces false positives).
- Large template:
 - less precise localization.
 - more robust to noise (good detector).
 - higher computational cost ($O(N \times N)$).

ASSIGNMENT 2: Playing with Sobel derivatives

Now that we have acquired a basic understanding of these methods, let's complete the following code cell to employ the Sobel kernels (S_x , S_y) to compute both derivatives and display them along with the original image (`medical_3.jpg`).

*Notice that the derivative image values can be positive **and negative**, caused by the negative values in the kernel. This implies that the desired depth of the destination image (`ddepth`) has to be at least a signed data type when calling to the `filter2D()` method.*

```

In [30]: # ASSIGNMENT 2
# Read one of the images, compute both kernel derivatives, apply the
# Write your code here!

# Read the image
image = cv2.imread(images_path + 'medical_3.jpg', 0)

# Define horizontal and vertical kernels
kernel_h = general_h_kernel(2)
kernel_v = general_v_kernel(2)


# Apply convolution
d_horizontal = cv2.filter2D(image, cv2.CV_16S, kernel_h) # Using ddepth
d_vertical = cv2.filter2D(image, cv2.CV_16S, kernel_v)

# And show them!
plt.subplot(131)
  
```

```
plt.imshow(image, cmap='gray')
plt.title('Original image')

plt.subplot(132)
plt.imshow(d_horizontal, cmap='gray')
plt.title('Horizontal derivative')

plt.subplot(133)
plt.imshow(d_vertical, cmap='gray')
plt.title('Vertical derivative');
```



Once we have computed both derivative images G_C and G_R , we can determine the *complete* edge image by computing the image gradient magnitude and then binarizing the result. Recall that the image codifying the gradient magnitude can be computed as:

$$|\nabla F(x, y)| = \sqrt{(F \otimes G_C)^2 + (F \otimes G_R)^2} \approx |F \otimes G_C| + |F \otimes G_R|$$

ASSIGNMENT 3a: Time to detect edges

Complete `edge_detection_chart()` that computes the gradient image of an input one using `kernel_h` and `kernel_v` (kernels for horizontal and vertical derivatives respectively) and **binarize the resultant image** (final edges image) using `threshold`. Then display in a 1x3 plot `image`, the gradient image, and finally, an image with the detected edges! (Only if `verbose` is True).

Tip: you should [normalize](https://docs.opencv.org/2.4/modules/core/doc/operations_on_arrays.html#normalize) (https://docs.opencv.org/2.4/modules/core/doc/operations_on_arrays.html#normalize) gradient image before thresholding.

Interesting functions: `np.absolute()` (<https://numpy.org/doc/stable/reference/generated/numpy.absolute.html>), `np.add()` (<https://numpy.org/doc/stable/reference/generated/numpy.add.html>), `cv2.threshold()` (https://docs.opencv.org/master/d7/d1b/group_imgproc_misc.html#gae8a4a146d1ca78c626a53577199e9c57)

```
In [31]: # ASSIGNMENT 3a
# Implement a function that that computes the gradient of an image,
# It must also binarize the resulting image using a threshold
# Show the input image, the gradient image (normalized) and the bin
def edge_detection_chart(image, kernel_h, kernel_v, threshold, verbose
    """ Computed the gradient of the image, binarizes and display it

    Args:
        image: Input image
        kernel_h: kernel for horizontal derivative
        kernel_v: kernel for vertical derivative
        threshold: threshold value for binarization
        verbose: Only show images if this is True
```



```

        Returns:
            edges: edges binary image
    """
    # Write your code here!

    # Compute derivatives
    d_h = cv2.filter2D(image, cv2.CV_16S, kernel_h) # horizontal
    d_v = cv2.filter2D(image, cv2.CV_16S, kernel_v) # vertical

    # Compute gradient
    gradient_image = np.absolute(np.add(d_h, d_v)) # Hint: You have

    # Normalize gradient
    norm_gradient = np.copy(image)
    norm_gradient = cv2.normalize(gradient_image, norm_gradient, 0,

    # Threshold to get edges
    ret, edges = cv2.threshold(norm_gradient, threshold, 255, cv2.THRESH_BINARY)

    if verbose:
        # Show the initial image
        plt.subplot(131)
        plt.imshow(image, cmap='gray')
        plt.title('Original image')

        # Show the gradient image
        plt.subplot(132)
        plt.imshow(gradient_image, cmap='gray')
        plt.title('Gradient image')

        # Show edges image
        plt.subplot(133)
        plt.imshow(edges, cmap='gray')
        plt.title('Edges detected')

    return edges

```

You can use next code to **test if your results are correct**:

In [32]: `image = np.array([[10,60,20],[60,22,74],[72,132,2]], dtype=np.uint8)`

```

# Sobel derivatives
kernel_h = np.array([[1,0,-1],[2,0,-2],[1,0,-1]])*1/4
kernel_v = np.array([[-1,-2,-1],[0,0,0],[1,2,1]])*1/4

print(edge_detection_chart(image, kernel_h, kernel_v, 100))

```

```

[[ 0  0  0]
 [255 255 255]
 [ 0 255  0]]

```

Expected output:

```

[[ 0  0  0]
 [255 255 255]
 [ 0 255  0]]

```

ASSIGNMENT 3b: Testing our detector

Now **try the implemented method** with different size Sobel kernels (3×3 , 5×5 , ...) and with other operators (Roberts, Prewitt, ...).

```
In [33]: # ASSIGNMENT 3b
# Read the image, set you kernels (Sobel, Roberts, Prewitt, etc.) and
# Write your code here!

# Read image
image = cv2.imread(images_path + 'medical_3.jpg', cv2.IMREAD_GRAYSCALE)

# Define kernel (Sobel, Roberts, Prewitt, ...)
# # Roberts
kernel_h = np.array([
    [0,0,0],
    [0,0,1],
    [0,-1,0]
])
kernel_v = np.array([
    [-1,0,0],
    [0,1,0],
    [0,0,0]
])

# Prewitt
# kernel_h = general_h_kernel(1)
# kernel_v = general_v_kernel(1)

# # Sobel 3x3
# kernel_h = general_h_kernel(2)
# kernel_v = general_v_kernel(2)

# Sobel 5x5
kernel_h = sobel_h_55()
kernel_v = sobel_v_55()

# # Frei-Chen
# kernel_h = general_h_kernel(np.sqrt(2))
# kernel_v = general_v_kernel(np.sqrt(2))

#Interact with your code!
interactive( edge_detection_chart, image=fixed(image), kernel_h=fixed(kernel_h), kernel_v=fixed(kernel_v), threshold=IntSlider(value=120, description='threshold', max=255, step=10), Output(), _dom_classes=...
```

Thinking about it (1)

Now, **answer following questions**:

- What happens if we use a bigger kernel?

Si usamos un kernel mas grande este puede dar errores de multiple respuesta ya que segun la imagenes donde se detectan los bordes, estos se ven más gruesos que un pixel, haciendo que no tenga tanta precision como un kernel pequeño

- There are differences between Sobel and other operators?

Las principales diferencias que dependen de la K en los kernels genericos como freichen, sobel y prewit la cual varian su deteccion del borde siendo más gruesos o finos

- What errors appear using those operators?


Con cualquiera de los operadores vamos a obtener errores de multiple respuesta y falsos negativos y positivos pero estos ultimos en menor medida si son kernels pequeños

- Why kernels usually are divided by a number? (e.g. 3×3 Sobel is divided by 4)

Para normalizar el kernel y que la operacion de convolucion los valores queden entre 0-255

3.1.2 DroG operator

Despite the simplicity of the previous techniques, they have a remarkable drawback: their performance is highly influenced by image noise. Taking a look at the following figure we can see how, having an apparently not so noisy function (first row), where it is easy to visually detect a step (an abrupt change in its values) around 1000, the response of the derivative with that level of noise is as bigger as the step itself!



Source: S. Seitz

But not everything is lost! An already studied image processing technique can be used to mitigate such noise: **image smoothing**, and more concretely, **Gaussian filtering**! The basic idea is to smooth the image and then apply a gradient operator, that is to compute $\frac{\partial}{\partial x}(f \otimes g)$. Not only that, this can be done even more efficiently thanks to the convolution derivative property:

$$\frac{\partial}{\partial x}(f \otimes g) = f \otimes \frac{\partial}{\partial x}g$$

That is, precomputing the resultant kernels from the convolution of the Gaussian filtering and the Sobel ones, and then convolving them with the image to be processed. With that we save one operation!

This combination of smoothing and gradient is usually called **Derivative of Gaussian operator (DroG)**. Formally:

$$\nabla[f(x, y) \otimes g_\sigma(x, y)] = f(x, y) \otimes \nabla[g_\sigma(x, y)] = f(x, y) \otimes \text{DroG}(x, y)$$

$$\text{DroG}(x, y) = \nabla [g_\sigma(x, y)] = \underbrace{\begin{bmatrix} \frac{\partial}{\partial x} [g_\sigma(x)g_\sigma(y)] \\ \frac{\partial}{\partial y} [g_\sigma(x)g_\sigma(y)] \end{bmatrix}}_{\text{separability}} = \underbrace{\begin{bmatrix} \frac{-xg_\sigma(x,y)}{\sigma^2} \\ \frac{-yg_\sigma(x,y)}{\sigma^2} \end{bmatrix}}_{g(x)' = -xg(x)/\sigma^2}$$

Remember from the previous notebooks the expression of the Gaussian distribution with 2 variables centered at the origin of coordinates, where the standard deviation σ controls the degree of smoothness:

$$g_\sigma(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

Take into account that the DroG template or kernel is **created just once!** Then it can be applied to as many images as you want.

ASSIGNMENT 4: Applying DroG

We would like to try this robust edge detection technique, so complete the `gaussian_kernel()` method that:

1. constructs a 2D gaussian filter (that is, $g_\sigma(x, y)$ in the previous DroG definition) from a 1D one, and
2. derives it, getting the DroG template (in other words, compute $-xg_\sigma(x, y)/\sigma^2$ and $-yg_\sigma(x, y)/\sigma^2$).
3. Finally, it calls our function `edge_detection_chart()`, but using the DroG template instead of the Sobel one.

Its inputs are:

- an image to be processed,
- the kernel aperture size,
- the standard deviation, and
- the gradient image binarization threshold.

```
In [23]: # ASSIGNMENT 4
# Implement a function that builds the horizontal and vertical DroG
```

```

# Inputs: an image, the kernel aperture size, the Gaussian standard deviation
# It returns the horizontal and vertical kernels
def gaussian_kernel(image, w_kernel, sigma, threshold, verbose=False):
    """ Construct the DroG operator and call edge_detection_chart.

    Args:
        image: Input image
        w_kernel: Kernel aperture size
        sigma: Standard deviation of the Gaussian distribution
        threshold: Threshold value for binarization
        verbose: Only show images if this is True

    Returns:
        DroG_h, DroG_v: DroG kernel for computing horizontal and vertical edges
    """
    # Write your code here!

    # Create the 1D gaussian filter
    s = sigma
    w = w_kernel
    gaussian_kernel_1D = np.array([(np.exp(-(z**2)/(2*(s**2))))/(s * np.sqrt(2*np.pi)) for z in range(-w, w+1)])

    # Get the 2D gaussian filter from the 1D one.
    vertical_kernel = gaussian_kernel_1D.reshape(2*w+1,1)
    horizontal_kernel = gaussian_kernel_1D.reshape(1,2*w+1)
    gaussian_kernel_2D = signal.convolve2d(vertical_kernel, horizontal_kernel, mode='same')

    # Construct DroG

    # Define x and y axis
    x = np.arange(-w,w+1)
    y = np.vstack(x)

    # Get the kernels for detecting horizontal and vertical edges
    DroG_h = x*(-gaussian_kernel_2D)/s**2 # Horizontal derivative
    DroG_v = y*(-gaussian_kernel_2D)/s**2 # Vertical derivative

    # Call edge detection chart using DroG
    edge_detection_chart(image, DroG_h, DroG_v, threshold, verbose)

    return DroG_h, DroG_v

```

You can use next code to **test if results are correct**:

```

In [24]: # Create an input image
image = np.array([[10,60,20],[60,22,74],[72,132,2]], dtype=np.uint8)

# Apply the Gaussian kernel
gaussian_kernel(image, w_kernel=1, sigma=1.2, threshold=100)

```

```

Out[24]: (array([[ 0.03832673, -0.05423735, -0.03832673],
                  [ 0.05423735, -0.03832673, -0.05423735],
                  [ 0.03832673, -0.05423735, -0.03832673]]),
          array([[ 0.03832673,  0.05423735,  0.03832673],
                  [-0.05423735, -0.03832673, -0.05423735],
                  [-0.03832673, -0.05423735, -0.03832673]]))

```

Expected output:

```
(array([[ 0.03832673, -0.          , -0.03832673],
       [ 0.05423735, -0.          , -0.05423735],
       [ 0.03832673, -0.          , -0.03832673]]),
array([[ 0.03832673,  0.05423735,  0.03832673],
       [-0.          , -0.          , -0.          ],
       [-0.03832673, -0.05423735, -0.03832673]]))
```

Thinking about it (2)

Now **try this method** and play with its interactive parameters in the next code cell. Then **answer the following questions**:

- What happens if a bigger kernel is used?

Pues que pierde detalle al suavizar, y por tanto no se detectan bordes donde deberian de estar

- What kind of errors appear and disappear whenever sigma is modified?

Cuando el sigma es muy pequeño el error predominante es el de fallo de detección (falsos positivos y falsos negativos) Y cuando es grande el error mas comun es de multiple respuesta ya que los bordes comienzan a ser mas gruesos

- Why the gradient image have lower values than the one from the original image? *Tip: [image normalization \(https://stackoverflow.com/questions/40645985/opencv-python-normalize-image/42164670\)](https://stackoverflow.com/questions/40645985/opencv-python-normalize-image/42164670)*

Por el normalizado que aplica OpenCV a la imagen

- Now that you have tried different techniques, in your opinion, which is the best one for this type of images?

Yo creo que la mejor tecnica seria la de DroG ya que con los parametros correctos podriamos eliminar la mayoria de errores con bastante eficacia

```
In [25]: # Read the image
image = cv2.imread(images_path + 'medical_3.jpg', 0)

# Interact with the three input parameters
interactive(gaussian_kernel, image=fixed(image), w_kernel=(1,5,1), s

interactive(children=(IntSlider(value=3, description='w_kernel', m
ax=5, min=1), FloatSlider(value=2.4, descrip..
```

ASSIGNMENT 5: Measuring efficiency

Finally, **you are asked to** compare the execution time of creating a DroG template using your previous code (combining two 1-D gaussian filters), with creating a DroG template by directly building the 2-D gaussian filter. Play with different (big) aperture sizes (up to 500).

Interesting resource: [how to measure execution time in Python \(https://stackoverflow.com/questions/14452145/how-to-measure-time-taken-between-lines-of-code-in-python\)](https://stackoverflow.com/questions/14452145/how-to-measure-time-taken-between-lines-of-code-in-python)

```

In [26]: # ASSIGNMENT 5
# Write your code here!

import time

print("Measuring the execution time needed for ...")

s = 1
w = 500

start = time.process_time() # Start timer

# FIRST WAY: Building the 2D gaussian filter from the 1D one

# Create 1D Gaussian filter
gaussian_kernel_1D = np.array([(np.exp(-(z**2)/(2*(s**2))))/(s * np

# Get the 2D gaussian filter from the 1D one.
vertical_kernel = gaussian_kernel_1D.reshape(2*w+1,1)
horizontal_kernel = gaussian_kernel_1D.reshape(1,2*w+1)
gaussian_kernel_2D = signal.convolve2d(vertical_kernel, horizontal_k

# Construct DroG

# Define x and y axis
x = np.arange(-w,w+1)
y = np.vstack(x)

# Get the kernels for detecting horizontal and vertical edges
DroG_h = gaussian_kernel_2D*(-x)/s**2 # Horizontal derivative
DroG_v = gaussian_kernel_2D*(-y)/s**2 # Vertical derivative
print("DroG building time using 1D Gaussian:", round(time.process_t

start = time.process_time() # Start timer

# SECOND WAY: Directly building the 2D gaussian filter

# Create 2D Gaussian filter
gaussian_kernel_2D = np.array([(np.exp(-((z**2)+(k**2))/(2*(s**2))

# Define x and y axis
x = np.arange(-w,w+1)
y = np.vstack(x)

# Construct DroG
DroG_h = gaussian_kernel_2D*(-x)/s**2 # Horizontal derivative
DroG_v = gaussian_kernel_2D*(-y)/s**2 # Vertical derivative

print("DroG building time using 2D Gaussian:", round(time.process_t

Measuring the execution time needed for ...
DroG building time using 1D Gaussian: 3.1673 seconds
DroG building time using 2D Gaussian: 3.78574 seconds

```

Conclusion


Awesome! Now you have expertise in more applications of the convolution operator. In this notebook you:

- Learned basic operators for edge detection that perform a **discrete approximation of a gradient operator**.
- Learned **how to construct a DroG kernel** in an efficient way.
- Played a bit with them in the context of medical images, discovering some real and meaningful utilities.

3.2 Operators based on second derivative


In the previous notebooks we saw how to detect edges by looking at the gradient image (first-derivative), but it is also possible to do that by analyzing the output of operators based on the second-derivative!

As you may remember, first derivative operators try to detect edges by looking for high magnitude values of such derivatives. The figure below shows a one-dimensional continuous function $f(x)$ in (a) and its first derivative in (b), where we can see that the point corresponding to the highest intensity difference reaches a maximum value:



The third figure (c) shows its second derivative, so we can check how such a value corresponds to... **a zero crossing!** That is, a second derivative yields a zero-crossing at points where the gradient presents a maximum, so we could detect edges looking for those crossings.

Unfortunately, things get a little tricky when moving to a 2D space (like images). Why? because depending on the orientation of the edge, this zero-crossing may go almost unnoticed (see, for example, d):




In this notebook we are going to explore two methods that face such issue and detect edges using the second derivative. These are:

- Laplacian operator ([Section 3.2.1](#))
- LoG operator ([Section 3.2.2](#))

Additionally, we will also take a look at a widely used algorithm that is a combination of different techniques: the **Canny algorithm** ([Section 3.2.3](#)).

Problem context - Edge detection for medical images

Unfortunately, you were not accepted (yet!) by the researching team at *Hospital Clínico* because the obtained results in the previous notebook were not as good as expected. Anyway, they have shown you the algorithms that they are currently using so you can study them for future opportunities. Let's have a look!



```
In [1]: import numpy as np
        from scipy import signal
        import cv2
        import matplotlib.pyplot as plt
        import matplotlib
        from ipywidgets import interact, fixed, widgets
        from mpl_toolkits.mplot3d import Axes3D

        matplotlib.rcParams['figure.figsize'] = (15.0, 15.0)

        images_path = './images/'
```

3.2.1 Laplacian operator

To face the previously posed issue about unnoticed edges due to the derivate orientation, the idea behind the Laplacian operator is to combine second derivatives in perpendicular directions. Thus, it is defined as:

$$\nabla^2 f(i, j) = \frac{\partial^2}{\partial x^2} f(i, j) + \frac{\partial^2}{\partial y^2} f(i, j)$$

Note that, by definition, **it returns a scalar**, not a vector as in the gradient case. Indeed, the Laplacian is the trace of the *Hessian matrix*, which fully characterizes the second derivative of a function:

$$H(f) = \begin{bmatrix} \frac{\partial^2 f}{\partial x^2} & \frac{\partial}{\partial x} \frac{\partial f}{\partial y} \\ \frac{\partial f}{\partial y} \frac{\partial f}{\partial x} & \frac{\partial^2 f}{\partial y^2} \end{bmatrix}$$

Compared with the first derivative-based edge detectors such as the Sobel operator, the

Laplacian operator have a number of advantages:

- it is a linear operator,
- invariant to image orientation, and
- precise when localizing edges.

Implementation

Now that we know the theory, let's have a look at how the Laplacian operator is implemented:

1. We start by considering first derivatives (OpenCV uses Sobel, but any alternative is valid):

$$\frac{\partial f(x, y)}{\partial x} = f_x(x, y) \approx G_R(i, j) = f(i + 1, j) - f(i, j)$$

$$\frac{\partial f(x, y)}{\partial y} = f_y(x, y) \approx G_C(i, j) = f(i, j + 1) - f(i, j)$$

2. Then, take second derivatives using the previous definitions:

$$g = \frac{\partial^2 f}{\partial x^2} = f_{xx}(x, y) \approx G_R(i, j) - G_R(i - 1, j) = f(i + 1, j) - 2f(i, j) + f(i - 1, j)$$

$$h = \frac{\partial^2 f}{\partial y^2} = f_{yy}(x, y) \approx G_C(i, j) - G_C(i, j - 1) = f(i, j + 1) - 2f(i, j) + f(i, j - 1)$$

3. Finally, implement it as a convolution with a certain kernel, so

$L[F(i, j)] = F(i, j) \otimes L(i, j)$. This would lead to the operation

$L[F(i, j)] = (F(i, j) \otimes g) + (F(i, j) \otimes h)$, but thanks to the distributive property of convolution:

$$\underbrace{f \otimes (g + h)}_{\text{One convolution}} = \underbrace{(f \otimes g) + (f \otimes h)}_{\text{Two convolutions}}$$

We can obtain a kernel that carries out both convolutions at once!:

g		h		g+h=L																											
<table border="1" style="border-collapse: collapse; width: 60px;"> <tr><td>0</td><td>0</td><td>0</td></tr> <tr><td>1</td><td>-2</td><td>1</td></tr> <tr><td>0</td><td>0</td><td>0</td></tr> </table>	0	0	0	1	-2	1	0	0	0	+	<table border="1" style="border-collapse: collapse; width: 60px;"> <tr><td>0</td><td>1</td><td>0</td></tr> <tr><td>0</td><td>-2</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>0</td></tr> </table>	0	1	0	0	-2	0	0	1	0	=	<table border="1" style="border-collapse: collapse; width: 60px;"> <tr><td>0</td><td>1</td><td>0</td></tr> <tr><td>1</td><td>-4</td><td>1</td></tr> <tr><td>0</td><td>1</td><td>0</td></tr> </table>	0	1	0	1	-4	1	0	1	0
0	0	0																													
1	-2	1																													
0	0	0																													
0	1	0																													
0	-2	0																													
0	1	0																													
0	1	0																													
1	-4	1																													
0	1	0																													

Zero-crossing

Note that the result of applying the Laplacian operator is not directly an edges image, but a second-derivative image. Recall that in the case of operators based on the first derivative we had to combine the two images returned by the gradient operator, and then apply a threshold to select edges. In this case, **it is needed an algorithm to detect**

zero-crossings in the second-derivative (Laplacian) image in order to return a binary image of edges.

An example of a simple zero-crossing algorithm could be:

1. Select a small positive number th (threshold).
2. A pixel is labelled as an edge if in the Laplacian image:
 - its value is smaller than $-th$ and at least one of its neighbours is bigger than th ,
or
 - its value is bigger than th and at least one of its neighbours is smaller than $-th$.

Advantages:

- Zero crossing produces a closed (or almost closed) contour, and
- it provides edges of 1-pixel width!

Limitations

- Unfortunately, the Laplacian operator is very sensitive to noise, resulting in a poor edge detection. Solution: If the image is blurred using a Gaussian filter before applying the Laplace operator, we can partially solve the noise problem. If this is done, the resultant is called **LoG (Laplacian of Gaussian)**.

3.2.2 LoG operator

So, the LoG operator first smoothes the image, and then applies the Laplacian operator (or viceversa, it's commutative!). Considering the convolution properties:

$$\nabla^2[f(x, y) \otimes g_\sigma(x, y)] = f(x, y) \otimes \nabla^2[g_\sigma(x, y)] = f(x, y) \otimes LoG_\sigma(x, y)$$

LoG is an isotropic operator, that is, it keeps radial symmetry. In this way, it is assumed that the covariance in both image dimensions is the same! Mathematically it is expressed as:

$$LoG_\sigma(x, y) = \frac{1}{\pi\sigma^4} \left[\frac{x^2 + y^2}{2\sigma^2} - 1 \right] \exp^{-\frac{x^2+y^2}{2\sigma^2}} = \frac{1}{\pi\sigma^4} \left[\frac{r^2}{2\sigma^2} - 1 \right] \exp^{-\frac{r^2}{2\sigma^2}} = LoG_\sigma(r')$$


Let's print the LoG operator!

```
In [2]: # Gauss filter
v = np.arange(-5,5,0.1)
X, Y = np.meshgrid(v,v)
covar = np.array([[2, 0],[0, 2]]) ## Assuming no correlation between
gauss_filter = np.exp(-0.5*(X**2/covar[0][0]+Y**2/covar[1][1]))

# Laplace filter
laplace_filter = np.array([[0,1,0],[1,-4,1],[0,1,0]], dtype="float")

# LoG operator
LoG = cv2.filter2D(gauss_filter, -1, laplace_filter)

# Plot it!
fig = plt.figure()
ax = fig.gca(projection='3d')
ax.plot_surface(X,Y,LoG,cmap='summer', edgecolor='none');
```



As a side note, the LoG operator is not separable. However, it can be implemented as **DoG (Difference of Gaussians)**, a sum of separable operators, reducing its complexity from $O(N^2)$ to $O(4N)$. The DoG is defined as:

$$DoG_{\sigma_1\sigma_2}(x, y) = g_{\sigma_1}(x, y) - g_{\sigma_2}(x, y) = g_{\sigma_1}(x)g_{\sigma_1}(y) - g_{\sigma_2}(x)g_{\sigma_2}(y)$$

Giving the ratio $\sigma_1/\sigma_2 = 1.6$ the best approximation of LoG. This complexity reduction

approach is employed, for example, in the popular SIFT keypoint detector, as we will see in following notebooks.

Limitations

- It is computationally costly,
- it doesn't provide any information about edge orientations,
- the output contains negative and non-integer values, so for display purposes the image should be normalized to the range 0-255,
- it is needed a zero-crossing method, and
- it tends to round object corners (more heavily as σ grows).

Experiencing Laplacian and LoG operators

Now that we are almost experts in the Laplacian and LoG operators, let's play a bit with them!

ASSIGNMENT 1a: Applying Gaussian smoothing

First, complete the function `gaussian_smoothing()` that:

1. blurs an image using a Gaussian filter, then
2. normalizes it to leverage the full range of values `[0, ..., 255]` (this is just a way to process the image in order to increase its contrast), and
3. finally returns the resulting image.

Interesting functions:

- For normalization you can use `cv2.normalize()` (https://docs.opencv.org/3.4/d2/de8/group_core_array.html#ga87eef7ee3970f86906d69a92cbf064bd).

```
In [3]: # ASSIGNMENT 1a
# Implement a function that blurs an input image using a Gaussian
def gaussian_smoothing(image, sigma, w_kernel):
    """ Blur and normalize input image.

    Args:
        image: Input image to be binarized
        sigma: Standard deviation of the Gaussian distribution
        w_kernel: Kernel aperture size

    Returns:
        smoothed_norm: Blurred image
    """
    # Write your code here!

    # Define 1D kernel
    s=sigma
    w=w_kernel
    kernel_1D = np.array([(np.exp(-(z**2)/(2*(s**2))))/(s * np.sqrt(2*np.pi)) for z in range(-w, w+1)])

    # Apply distributive property of convolution
    vertical_kernel = kernel_1D.reshape(2*w+1,1)
    horizontal_kernel = kernel_1D.reshape(1,2*w+1)
```

```

gaussian_kernel_2D = signal.convolve2d(vertical_kernel, horizontal_kernel, mode='same')

# Blur image
smoothed_img = cv2.filter2D(image, cv2.CV_8U, gaussian_kernel_2D)

# Normalize to [0 255] values
smoothed_norm = np.array(image.shape)
smoothed_norm = cv2.normalize(smoothed_img, None, 0, 255, cv2.NORM_MINMAX)

return smoothed_norm

```

ASSIGNMENT 1b: Detecting edges with Laplacian and LoG

Now, we are going to see the differences between the Laplacian and LoG operators. For that complete the `laplace_testing()` function which:

1. applies the Laplacian operator to the input image and
2. to a blurred version of the input image (use the previously implemented function `gaussian_smoothing()` to smooth it). Notice that applying the Laplacian operator after smoothing the image is equivalent to applying the LoG operator.
3. Finally displays both images along with the original one in a 1x3 plot.

This function uses as inputs:

- an image to be processed,
- the size of the Laplacian filter (should be odd), and
- the parameters of the Gaussian filter.

Note that it would be possible to reduce the computation time by precomputing LoG (as commented above). This is convolving the Laplacian and Gaussian filters instead of applying them separately.

Interesting functions:

- OpenCV defines the Laplace operator as `cv2.Laplacian()` (https://docs.opencv.org/3.4/d5/db5/tutorial_laplace_operator.html).

```

In [4]: # ASSIGNMENT 1b
# Implement a function that applies the Laplacian operator to the input image
# Display a 1x3 plot with the original image and the two resulting images
# Inputs: image, size of the Laplacian kernel, sigma and size of the Gaussian kernel
def laplace_testing(image, size_Laplacian, sigma, w_gaussian):
    """ Apply Laplacian and Log operators to an image.

    Args:
        image: Input image to be binarized
        size_Laplacian: size of Laplacian kernel (odd)
        sigma: Standard deviation of the Gaussian distribution
        w_gaussian: Gaussian kernel aperture size

    """
    # Write your code here!

    # Blur image
    blurred_img = gaussian_smoothing(image, sigma, w_gaussian)

    # Apply Laplacian to the original image

```

```

laplacian = cv2.Laplacian(image, cv2.CV_8U, ksize=size_Laplacian)

# Aplay Laplacian to the blurred image
laplacian_blurred = cv2.Laplacian(blurred_img, cv2.CV_8U, ksize=size_Laplacian)

# Show initial image
plt.subplot(131)
plt.imshow(image, cmap='gray')
plt.title('Original image')

# Show laplacian
plt.subplot(132)
plt.imshow(laplacian, cmap='gray')
plt.title('Laplacian without blurring')

# Show LoG
plt.subplot(133)
plt.imshow(laplacian_blurred, cmap='gray')
plt.title('Laplacian blurred (LoG)')

```

It is time to try this method to our medical images and play with interactive parameters.

```

In [5]: # Read an image
image = cv2.imread(images_path + 'medical_3.jpg', 0)

# Interact with the parameters!
interact(laplace_testing, image=fixed(image), size_Laplacian=(1,7,2

interactive(children=(IntSlider(value=3, description='size_Laplacian', max=7, min=1, step=2), FloatSlider(value=2, description='sigma', max=10, min=0.5, step=0.5)),

```

Thinking about it (1)

Now, **answer the following questions:**

- Could be the Laplacian applied without a previous blurring? Does this have any drawback?

Si pero habria mucho ruido y esto impediria que hubises una correcta deteccion de los bordes
- Are the images obtained in the previous function *edge images*?

No
- If not, what would be needed for obtaining the edges from those images?

Aplicarle un umbral para diferenciar correctamente los bordes

3.2.3 The Canny algorithm

The Canny edge detector^[1] is an algorithm that combines a number of techniques:



- the DroG operator,
- non-maxima suppression, and

- hysteresis.

It was designed to be a good detector, yield a good localization, and to provide a single response!

This algorithm consists of the following steps:


1. **Noise filtering and gradient image.** Apply the DroG operator to reduce noise and obtain a gradient image.
2. **Non-maximum suppression.** This removes pixels that are not considered to be part of an edge. Typically, the gradient image obtained after using DroG presents thick edges. The idea is to keep only those pixels that are maximum within their neighborhood in the direction of the gradient, suppressing the rest of them. Hence, only thin lines (candidate edges) will remain. For that:
 - We consider 4 main directions or *angular sectors*: $[0, 45]$, $[45, 90]$, $[90, 135]$, $[135, 180]$. The gradient angle $\theta[i, j]$ is approximated by where it lays.
 - A 3x3 filter is moved over the gradient image $G[i, j]$ at each pixel, and it suppresses the edge strength of the center pixel (for example by setting its value to 0) if its magnitude is not greater than the magnitude of the two neighbors in the gradient direction. This way we have a single response at each edge.



We pick the maximum of these three pixels

3. **Hysteresis:** The final step, for which the Canny algorithm uses two thresholds (upper and lower) to determine edge pixels:

- If the grey level of a candidate pixel of the gradient image is higher than the upper threshold, the pixel is accepted as an edge.
- If the grey level of a candidate pixel of the gradient image is below the lower threshold, then it is rejected.
- If the grey level of a candidate pixel of the gradient image is between the two thresholds, then it will be accepted only if it is connected to a pixel that is above the upper threshold and rejected otherwise.



This algorithm can be executed repeatedly with different levels of smoothing (changing the sigma of the DroG operator). Different sigmas produce edges at different spatial

ASSIGNMENT 2: The enormously popular Canny algorithm

Complete `canny_testing()`, which applies the Canny algorithm. Note that OpenCV Canny's implementation does not apply Gaussian smoothing, but directly applies Sobel. This gives to us the opportunity to:

1. check the performance of this technique by considering the initial image and a smoothed version of it. *Note: use our popular `gaussian_smoothing()` function for blurring the image*
2. After this, display both resulting images along the original one.

This function takes as arguments:

- an image,
- both lower and upper Canny thresholds, and
- the parameters of the Gaussian filter.

Interesting functions:

- OpenCV implements the Canny algorithm in `cv2.Canny()` (https://docs.opencv.org/2.4/modules/imgproc/doc/feature_detection.html?highlight=canny).

```
In [8]: # ASSIGNMENT 2
# Implement a function that applies the Canny operator to an input
# Display a 1x3 plot with the original image and the two resulting
# Inputs: image, size of the Laplacian kernel, sigma and size of the
def canny_testing(image, lower_threshold, upper_threshold, sigma, w_gaussian):
    """ Apply Canny algorithm to an image.

    Args:
        image: Input image to be binarized
        lower_threshold: bottom value for hysteresis
        upper_threshold: top value for hysteresis
        sigma: Standard deviation of the Gaussian distribution
        w_gaussian: Gaussian kernel aperture size

    """

    # Smooth image
    blurred_img = gaussian_smoothing(image, sigma, w_gaussian)

    # Apply Canny to original image
    canny = cv2.Canny(image, lower_threshold, upper_threshold)

    # Apply Canny to blurred image
    canny_blurred = cv2.Canny(blurred_img, lower_threshold, upper_thro

    # Show initial image
    plt.subplot(131)
    plt.imshow(image, cmap='gray')
    plt.title('Original image')

    # Show Canny without blurring
```

```
plt.subplot(132)
plt.imshow(canny, cmap='gray')
plt.title('Canny without smoothing')

# Show Canny with blurring
plt.subplot(133)
plt.imshow(canny_blurred, cmap='gray')
plt.title('Canny smoothed')
```

Among the multiple parameters of this algorithm, it is interesting to check its performance with different levels of smoothing (changing the sigma of the DroG operator). As commented, different sigma produces edges at different spatial features. **Try the effect of this and other parameters** playing with the interactive parameters in the following code cell. You can also try with your own images.

```
In [9]: # Read an image
image = cv2.imread(images_path + 'medical_2.jpg', 0)

# Interact with the parameters
interact(canny_testing, image=fixed(image), lower_threshold=(0,260,10),
upper_threshold=(10,260,10))

interactive(children=(IntSlider(value=120, description='lower_threshold', max=260, step=20), IntSlider(value=10, description='upper_threshold', max=260, step=20)),
```

Thinking about it (2)

Now, **answer following questions:**

- Could Canny be applied without a previous blurring? Which are the consequences of this?

Si, pero este tendria mucho ruido y no se podria diferenciar bien los bordes

- What is a *good* value for both, lower and upper thresholds? Would these values be the same for any input image?

Unos buenos valores para delimitar solamente el hueso serian (60, 160). No habrian que reajustarlos segun la imagen y lo que se quiera buscar en ella

- Now that you have tried a good number of edge detection methods, **which one is your favorite, and why?**

Tras haber visto todos los operadores de este capitulo, podria decir que si tenemos potencia de computo optaria por Canny

Conclusion

Terrific! You finished this notebook, that includes information about:

- Laplacian and LoG operators and the importance of smoothing, and
- how the Canny algorithm is implemented and how to use it.

Curiosity

The Canny algorithm is a well known algorithm in the computer vision field. It is used in a lot of modern technologies. However, the original paper was published in 1986 by John

References

[1]: CANNY, John. [A computational approach to edge detection.](https://www.google.com/url?sa=t&rct=j&q=&esrc=s&source=web&cd=10&ved=2ahUKEwiU9uyiganoAhWNDWMBHducCvsQFjAJegQIBhAB&url=http%3A%2F%2Fciteseerx.ist.psu.edu%2Fviewdoc%2Fdownload%3Fdoi%3D10.1.1.420.3300%26rep%3Dusg=AOvVaw3tsKoxnc3qnS7bji3HmvQc) (https://www.google.com/url?sa=t&rct=j&q=&esrc=s&source=web&cd=10&ved=2ahUKEwiU9uyiganoAhWNDWMBHducCvsQFjAJegQIBhAB&url=http%3A%2F%2Fciteseerx.ist.psu.edu%2Fviewdoc%2Fdownload%3Fdoi%3D10.1.1.420.3300%26rep%3Dusg=AOvVaw3tsKoxnc3qnS7bji3HmvQc). IEEE Transactions on pattern analysis and machine intelligence, 1986, no 6, p. 679-698.